Data Pre-processing

Zhuocheng Lin

3/18/2020

Packages

```
library(tidyverse)
library(modelr)
library(lubridate)
library(caret)
```

Read data

```
df <- read_csv('./US_Accidents_Dec19.csv', col_types = cols(.default = col_character())) %>% type_conve
```

Drop variables with high NA proportion (over 50%)

```
df %>% summarise_all(~ mean(is.na(.))) %>%
  pivot_longer(1:49, names_to = "variable", values_to = "NA_prop") %>%
  filter(NA_prop >= 0.5)
```

```
## # A tibble: 5 x 2
##   variable          NA_prop
##   <chr>               <dbl>
## 1 End_Lat             0.755
## 2 End_Lng             0.755
## 3 Number              0.645
## 4 Wind_Chill(F)       0.623
## 5 Precipitation(in)   0.672
```

```
drop_na_cols <- c("End_Lat", "End_Lng", "Number", "Wind_Chill(F)", "Precipitation(in)")
```

Drop unuseful variable

```
# these variables are not useful in predicting Severity
not_useful <- c("ID", "Source", "Timezone", "Airport_Code", "Weather_Timestamp", "Wind_Direction", "Cou

# Not so sure about whether Wind_Direction is useful
# to me, the relation seems weak
df %>% ggplot(aes(Wind_Direction, ..prop..)) +
    geom_bar(aes(group = Severity, fill = Severity)) +
```

```
    coord_flip() +
    labs(title = "Wind_Direction distribution in each level")
```

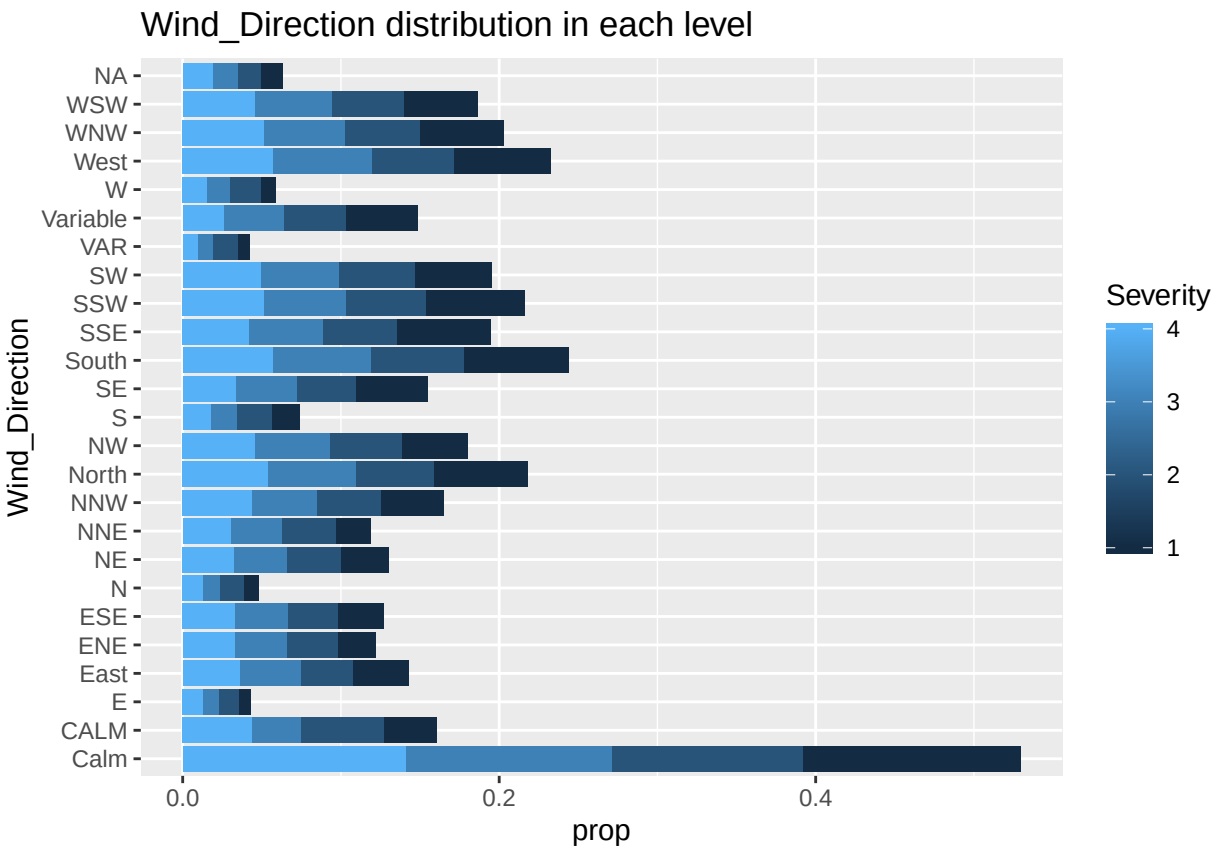


```
df_drop <- df %>% select(-drop_na_cols, -not_useful)
```

Rename variables to avoid potential error

```
df_drop <-  df_drop %>%
  rename("Distance" = `Distance(mi)`, "Temperature" = `Temperature(F)`, "Humidity" = `Humidity(%)`,
         "Pressure" = `Pressure(in)`, "Visibility" = `Visibility(mi)`, "Wind_Speed" = `Wind_Speed(mph)`)
```

Pre-processing time related variables

```
df_time <- df_drop %>%
  mutate(Duration = as.numeric(End_Time - Start_Time)) %>%
  # accident duration should be positive
  filter(!(Duration < 0)) %>%
  separate(Start_Time, into = c("Date", "Time"), sep = " ") %>%
  mutate("Year" = str_sub(Date, 1, 4), "Month" = str_sub(Date, 6, 7), "Day" = str_sub(Date, 9, 10), "Wda
  select(-c("Date", "Time", "End_Time")) %>%
  select(TMC, Severity, Year, Month, Day, Hour, Wday, Duration, everything())
head(df_time)
```

```
## # A tibble: 6 x 40
```

```
##     TMC Severity Year  Month Day   Hour  Wday  Duration Start_Lat Start_Lng
##   <dbl>    <dbl> <chr> <chr> <chr> <chr> <chr>    <dbl>     <dbl>     <dbl>
## 1   201        3 2016  02    08    05    2        18840      39.9     -84.1
## 2   201        2 2016  02    08    06    2         1800      39.9     -82.8
## 3   201        2 2016  02    08    06    2         1800      39.1     -84.0
## 4   201        3 2016  02    08    07    2         1800      39.7     -84.2
## 5   201        2 2016  02    08    07    2         1800      39.6     -84.2
## 6   201        3 2016  02    08    07    2         1800      40.1     -82.9
## # ... with 30 more variables: Distance <dbl>, Street <chr>, Side <chr>,
## #   City <chr>, County <chr>, State <chr>, Zipcode <chr>, Temperature <dbl>,
## #   Humidity <dbl>, Pressure <dbl>, Visibility <dbl>, Wind_Speed <dbl>,
## #   Weather_Condition <chr>, Amenity <lgl>, Bump <lgl>, Crossing <lgl>,
## #   Give_Way <lgl>, Junction <lgl>, No_Exit <lgl>, Railway <lgl>,
## #   Roundabout <lgl>, Station <lgl>, Stop <lgl>, Traffic_Calming <lgl>,
## #   Traffic_Signal <lgl>, Turning_Loop <lgl>, Sunrise_Sunset <chr>,
## #   Civil_Twilight <chr>, Nautical_Twilight <chr>, Astronomical_Twilight <chr>
```

Address

```
# not sure the best way to deal with address
# my opinion is we can choose one state data, and build the model
# and ignore Street, County and City
address <- c("Street", "County", "City", "Zipcode")
df_add <- df_time %>% select(-address)
```

Drop missing Weather_Condition

```
# when Weather_Condition is missing,
# other variables related to weather will be missing too (most cases)
df_add %>% filter(is.na(Weather_Condition)) %>% select(Temperature:Weather_Condition)
```

```
## # A tibble: 65,932 x 6
##    Temperature Humidity Pressure Visibility Wind_Speed Weather_Condition
##          <dbl>    <dbl>    <dbl>      <dbl>      <dbl> <chr>
##  1        48.2       93     29.5         10        9.2 <NA>
##  2        NA         NA     NA           NA       NA   <NA>
##  3        95         20     29.9         10        6.9 <NA>
##  4        91.4       28     29.9         10       15   <NA>
##  5        NA         NA     NA           NA       NA   <NA>
##  6        NA         NA     NA           NA       NA   <NA>
##  7        NA         NA     NA           NA       NA   <NA>
##  8        NA         NA     NA           NA       NA   <NA>
##  9        NA         NA     NA           NA       NA   <NA>
## 10        NA         NA     NA           NA       NA   <NA>
## # ... with 65,922 more rows
```

```
df_add %>% filter(is.na(Weather_Condition)) %>% select(Temperature:Weather_Condition) %>% summarise_all
```

```
## # A tibble: 1 x 6
##   Temperature Humidity Pressure Visibility Wind_Speed Weather_Condition
##         <int>    <int>    <int>      <int>      <int>             <int>
## 1       46246    46309    44532      58500      56084             65932
```

```
# we can drop observations whose Weather_Condition is missing
df_weather <- df_add %>% filter(!is.na(Weather_Condition))
```

Format

```
df_weather <- df_weather %>%
  mutate(TMC = as.character(TMC)) %>%
  mutate_if(is.logical, as.character)
```

Replace NA with mean

```
df_mean <- df_weather %>%
  mutate_if(is.numeric, ~ replace_na(., mean(., na.rm = T)))
summary(df_mean)
```

```
##      TMC                Severity         Year              Month
##  Length:2908381     Min.   :1.000   Length:2908381     Length:2908381
##  Class :character   1st Qu.:2.000   Class :character   Class :character
##  Mode  :character   Median :2.000   Mode  :character   Mode  :character
##                     Mean   :2.359
##                     3rd Qu.:3.000
##                     Max.   :4.000
##      Day                Hour               Wday              Duration
##  Length:2908381     Length:2908381     Length:2908381     Min.   :      73
##  Class :character   Class :character   Class :character   1st Qu.:    1783
##  Mode  :character   Mode  :character   Mode  :character   Median :    2675
##                                                           Mean   :    7063
##                                                           3rd Qu.:    4481
##                                                           Max.   :91680802
##    Start_Lat       Start_Lng          Distance            Side
##  Min.   :24.56   Min.   :-124.62   Min.   :  0.0000   Length:2908381
##  1st Qu.:33.54   1st Qu.:-117.30   1st Qu.:  0.0000   Class :character
##  Median :35.82   Median : -90.25   Median :  0.0000   Mode  :character
##  Mean   :36.48   Mean   : -95.47   Mean   :  0.2831
##  3rd Qu.:40.41   3rd Qu.: -80.95   3rd Qu.:  0.0100
##  Max.   :49.00   Max.   : -67.11   Max.   :333.6300
##     State            Temperature        Humidity         Pressure
##  Length:2908381     Min.   :-40.00   Min.   :  1.00   Min.   : 0.00
##  Class :character   1st Qu.: 50.00   1st Qu.: 49.00   1st Qu.:29.82
##  Mode  :character   Median : 64.40   Median : 67.00   Median :29.98
##                     Mean   : 62.38   Mean   : 65.41   Mean   :29.83
##                     3rd Qu.: 76.00   3rd Qu.: 84.00   3rd Qu.:30.11
##                     Max.   :170.60   Max.   :100.00   Max.   :33.04
##    Visibility        Wind_Speed      Weather_Condition    Amenity
##  Min.   :  0.000   Min.   :  0.000   Length:2908381     Length:2908381
##  1st Qu.: 10.000   1st Qu.:  5.800   Class :character   Class :character
##  Median : 10.000   Median :  8.100   Mode  :character   Mode  :character
##  Mean   :  9.151   Mean   :  8.296
##  3rd Qu.: 10.000   3rd Qu.: 10.400
##  Max.   :140.000   Max.   :822.800
##      Bump            Crossing           Give_Way           Junction
```

```
##  Length:2908381     Length:2908381     Length:2908381     Length:2908381
##  Class :character   Class :character   Class :character   Class :character
##  Mode  :character   Mode  :character   Mode  :character   Mode  :character
##
##
##
##    No_Exit            Railway           Roundabout          Station
##  Length:2908381     Length:2908381     Length:2908381     Length:2908381
##  Class :character   Class :character   Class :character   Class :character
##  Mode  :character   Mode  :character   Mode  :character   Mode  :character
##
##
##
##      Stop           Traffic_Calming    Traffic_Signal     Turning_Loop
##  Length:2908381     Length:2908381     Length:2908381     Length:2908381
##  Class :character   Class :character   Class :character   Class :character
##  Mode  :character   Mode  :character   Mode  :character   Mode  :character
##
##
##
##  Sunrise_Sunset     Civil_Twilight     Nautical_Twilight  Astronomical_Twilight
##  Length:2908381     Length:2908381     Length:2908381     Length:2908381
##  Class :character   Class :character   Class :character   Class :character
##  Mode  :character   Mode  :character   Mode  :character   Mode  :character
##
##
##
```

TMC

```
# most TMC NAs are in level 4
df_mean %>%
  ggplot(aes(TMC, ..prop..)) +
    geom_bar(aes(group = Severity, fill = Severity)) +
    labs(title = "TMC distribution in each severity level")
```

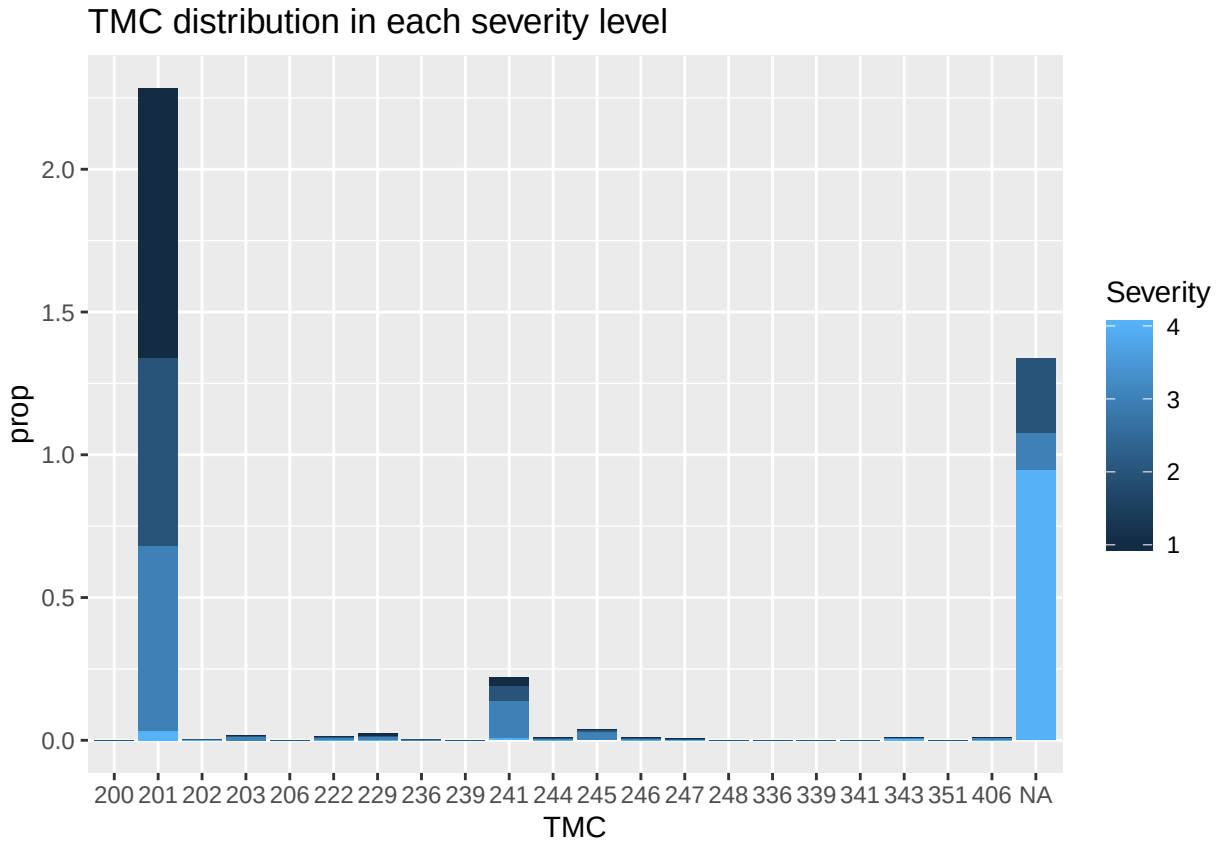


```
df_mean %>%
  ggplot(aes(TMC, ..prop..)) +
    geom_bar(aes(group = Severity, fill = Severity)) +
    labs(title = "TMC distribution in each severity level") +
    coord_cartesian(ylim = c(0, 0.25))
```

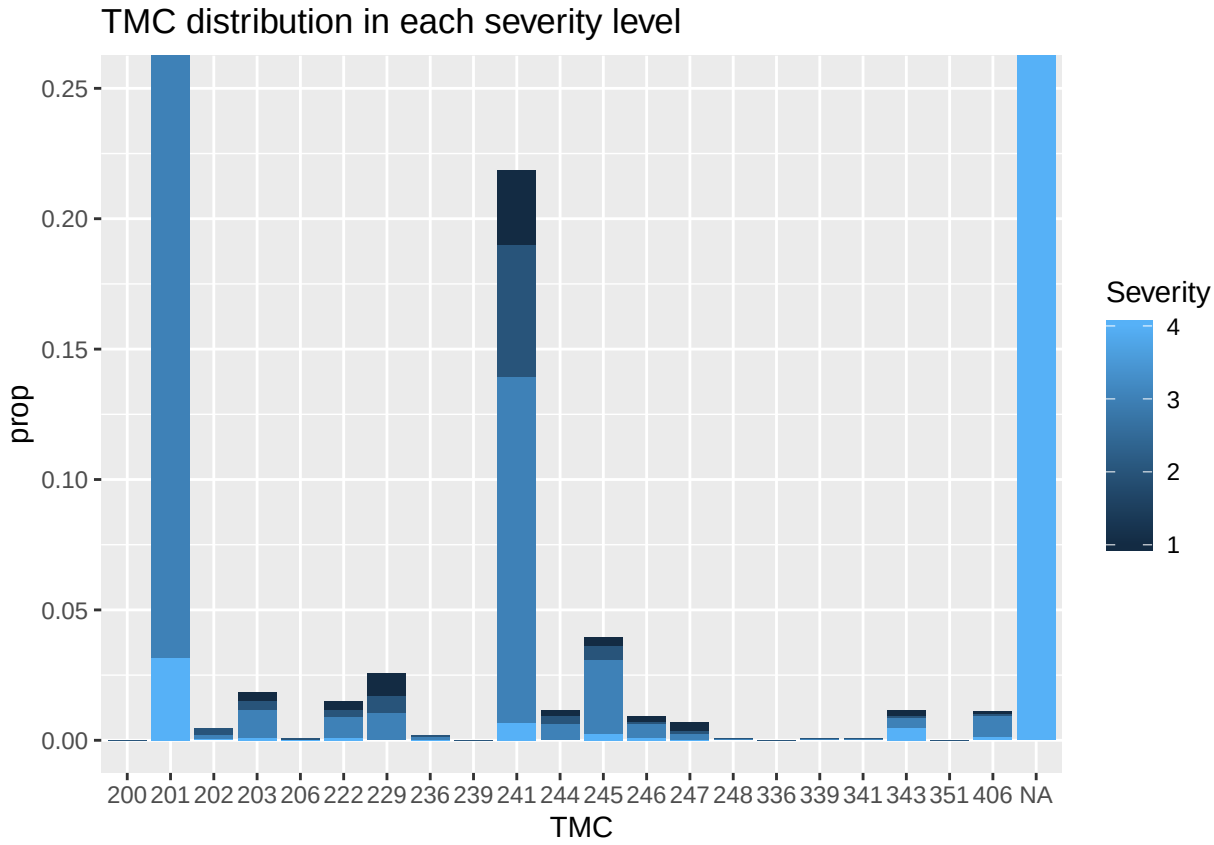


```
# my opinion is TMC NA can be considered as an important feature of Severity
# we can treate NA as a new TMC code
df_TMC <- df_mean %>%
  mutate(TMC = replace_na(TMC, "NA"))
```

Final check if there is unusual observation

```
df_TMC %>% summarise_all(~sum(is.na(.))) %>% pivot_longer(everything(), names_to = "variable", values_t
```

```
## # A tibble: 5 x 2
##   variable                NAs
##   <chr>                 <int>
## 1 Side                      1
## 2 Sunrise_Sunset           80
## 3 Civil_Twilight           80
## 4 Nautical_Twilight        80
## 5 Astronomical_Twilight    80
```

```
# Side has 1 NA, remove it
# variables related to daylight all have 80 NAs

df_TMC %>% filter(is.na(Sunrise_Sunset)) %>% count(TMC)
```

```
## # A tibble: 6 x 2
##   TMC       n
##   <chr> <int>
```

```
## 1 201      39
## 2 222       1
## 3 229       2
## 4 241       2
## 5 343       1
## 6 NA       35
```

```r
# the missing daylight data may be related to missing TMC
# replace them with a new levle "NAs"
df_final <- df_TMC %>%
  filter(!is.na(Side)) %>%
  mutate_at(vars(Sunrise_Sunset:Astronomical_Twilight), ~ replace_na(., "NA"))
```

Write csv file

```r
# write_csv(df_final, "./tidy.csv")
```