

Different Models

Zhuocheng Lin

3/22/2020

1. Packages

```
library(tidyverse)
library(modelr)
library(caret)
library(ROSE)
library(randomForest)
library(glmnet)
library(rpart)
```

2. Data pre-processing

(1) Read data

```
df <- read_csv("./tidy.csv", col_types = cols(.default = col_character())) %>%
  type_convert()
```

(2) Specify factors

```
df_format <- df %>%
  mutate(TMC = factor(TMC), Severity = factor(Severity), Year = factor(Year), Wday = factor(Wday)) %>%
  mutate_if(is.logical, factor) %>%
  mutate_if(is.character, factor)
```

(3) Narrow down to one State

```
df_format %>% count(State) %>% arrange(desc(n))
```

```
## # A tibble: 49 x 2
##   State      n
##   <fct> <int>
## 1 CA      650285
## 2 TX      291281
## 3 FL      221148
## 4 SC      143606
## 5 NC      141397
## 6 NY      136288
## 7 PA       89120
## 8 MI       88488
```

```
## 9 IL      86105
## 10 GA     82547
## # ... with 39 more rows
# choose TX as the target State
df_TX <- df_format %>% filter(State == "TX") %>% select(-State)
```

(4) Remove unuseful variables

```
# remove variables with only 1 distinct value
df_TX %>% summarise_all(~ n_distinct(.)) %>%
  pivot_longer(everything(), names_to = "variable", values_to = "n") %>% filter(n == 1)

## # A tibble: 1 x 2
##   variable      n
##   <chr>      <int>
## 1 Turning_Loop    1

df_TX <- df_TX %>% select(-Turning_Loop)
```

(5) Drop weather condition levels

```
# some Weather_Condition levels only have a few observations
# which can be a problem when we try to build a model
df_TX %>% count(Weather_Condition) %>% filter(n < 20) %>% select(Weather_Condition)

## # A tibble: 18 x 1
##   Weather_Condition
##   <fct>
## 1 Blowing Dust
## 2 Drizzle and Fog
## 3 Haze / Windy
## 4 Heavy Drizzle
## 5 Heavy T-Storm / Windy
## 6 Light Drizzle / Windy
## 7 Light Freezing Fog
## 8 Light Haze
## 9 Light Ice Pellets
## 10 Light Rain Showers
## 11 Light Snow / Windy
## 12 N/A Precipitation
## 13 Rain Showers
## 14 Sand
## 15 Showers in the Vicinity
## 16 Smoke
## 17 Thunder / Windy
## 18 Wintry Mix

drop_weather <- df_TX %>% count(Weather_Condition) %>% filter(n < 20) %>% select(Weather_Condition)
drop_weather <- drop_weather$Weather_Condition %>% unlist()
df_TX <- df_TX %>% filter(!(Weather_Condition %in% drop_weather))
df_TX <- df_TX %>% mutate(Weather_Condition = factor(Weather_Condition))
```

(6) Add new labels

```
# group level 3 and 4 together, as "Severe"
# group level 1 and 2 together, as "Not Severe"
df_label <- df_TX %>%
  mutate("Status" = factor(ifelse(Severity == "3" | Severity == "4", "Severe", "Not Severe"),
                             levels = c("Not Severe", "Severe")))
df_label %>% select(Severity, Status)

## # A tibble: 291,156 x 2
##   Severity Status
##   <fct>    <fct>
## 1 2      Not Severe
## 2 2      Not Severe
## 3 2      Not Severe
## 4 2      Not Severe
## 5 3      Severe
## 6 2      Not Severe
## 7 2      Not Severe
## 8 3      Severe
## 9 3      Severe
## 10 2     Not Severe
## # ... with 291,146 more rows
```

(7) Near Zero-Variance Predictors

```
# these variable may become zero-variance when the data are split into subsets
# remove them
nzv <- nearZeroVar(df_label, saveMetrics = T)
nzv[nzv$nzv,]
```

```
##           freqRatio percentUnique zeroVar  nzv
## Visibility      21.90956    0.015112174  FALSE TRUE
## Amenity         52.04354    0.000686917  FALSE TRUE
## Bump           5598.15385    0.000686917  FALSE TRUE
## Give_Way       176.10219    0.000686917  FALSE TRUE
## No_Exit        969.52000    0.000686917  FALSE TRUE
## Railway       106.71587    0.000686917  FALSE TRUE
## Roundabout    26467.72727    0.000686917  FALSE TRUE
## Station        62.64066    0.000686917  FALSE TRUE
## Stop          66.13304    0.000686917  FALSE TRUE
## Traffic_Calming 2598.60714    0.000686917  FALSE TRUE
```

```
nzv_cols <- rownames(nzv[nzv$nzv,])
df_label <- df_label %>%
  select(-nzv_cols)
```

(8) Partition

```
set.seed(1)
df_parts <- resample_partition(df_label, c(train = 0.6, valid = 0.2, test = 0.2))
train <- as_tibble(df_parts$train)
valid <- as_tibble(df_parts$valid)
test <- as_tibble(df_parts$test)
```

```

# check Weather_Condition levels
# train should have more levels than valid and test
tr <- train %>% select(Weather_Condition) %>% distinct()
va <- valid %>% select(Weather_Condition) %>% distinct()
te <- test %>% select(Weather_Condition) %>% distinct()
setdiff(va, tr)

```

```

## # A tibble: 0 x 1
## # ... with 1 variable: Weather_Condition <fct>

```

```

setdiff(te, tr)

```

```

## # A tibble: 0 x 1
## # ... with 1 variable: Weather_Condition <fct>

```

(9) Sampling

```

new_train <- ovun.sample(Status ~ .,
                          data = train %>% select(-Severity),
                          method = "both", p = 0.5, N = 90000)$data %>% as_tibble()
table(new_train$Status)

```

```

##
## Not Severe      Severe
##      44978      45022

```


3. Use different models to fit the data

(1) Sparse Logistic regression


```

x <- model.matrix(Status ~ ., data = new_train)
model_total <- glmnet(x, new_train$Status, family = "binomial")
plot(model_total, xvar = "lambda", label = T)

```



```
model_lambda <- cv.glmnet(x, new_train$Status, family = "binomial")
plot(model_lambda)
```



```
# use the best lambda
valid_pred <- valid %>%
  mutate("pred" = predict(model_lambda,
                           newx = model.matrix(Status ~ ., data = valid %>% select(-Severity)),
                           s = "lambda.min", type = "response")[,1]) %>%
  mutate("pred" = ifelse(pred > 0.5, "Severe", "Not Severe"))
valid_pred %>% select(Status, pred)
```

```
## # A tibble: 58,231 x 2
##   Status      pred
##   <fct>      <chr>
## 1 Not Severe Not Severe
## 2 Not Severe Not Severe
## 3 Not Severe Not Severe
## 4 Severe     Severe
## 5 Severe     Severe
## 6 Not Severe Not Severe
## 7 Not Severe Severe
## 8 Not Severe Not Severe
## 9 Not Severe Not Severe
## 10 Not Severe Not Severe
## # ... with 58,221 more rows
```

```
table(valid$Status)
```

```
##
## Not Severe      Severe
```

```
##      42154      16077
confusionMatrix(table(valid_pred$Status, valid_pred$pred))


## Confusion Matrix and Statistics
##
##
##           Not Severe Severe
## Not Severe      31410 10744
## Severe          4332 11745
##
##           Accuracy : 0.7411
##           95% CI : (0.7375, 0.7447)
## No Information Rate : 0.6138
## P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.4234
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.8788
##           Specificity : 0.5223
##           Pos Pred Value : 0.7451
##           Neg Pred Value : 0.7305
##           Prevalence : 0.6138
##           Detection Rate : 0.5394
## Detection Prevalence : 0.7239
##           Balanced Accuracy : 0.7005
##
##           'Positive' Class : Not Severe
##
```

(2) Random forest

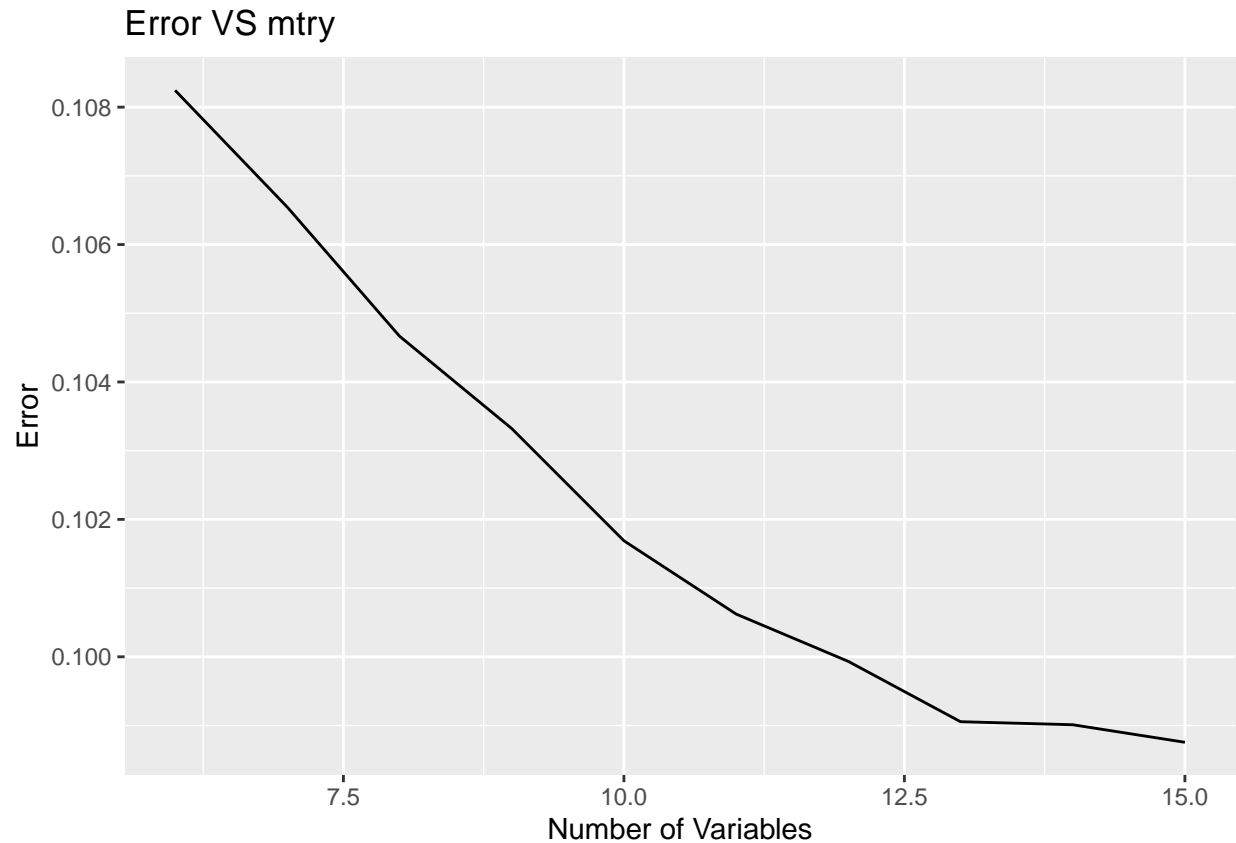
```
model <- randomForest(Status ~ ., data = new_train, mtry = 6, ntree = 500)

# see if ntree = 500 is enough
error_data <- model$err.rate %>%
  as_tibble() %>%
  mutate("Trees" = seq_along(OOB)) %>%
  pivot_longer(cols = 1:3, names_to = "Type", values_to = "Error")

ggplot(error_data, aes(Trees, Error, color = Type)) +
  geom_line() +
  labs(x = "Number of Trees",
       title = "Error Rate")
```



```
# try different mtry
oob_values <- vector(length = 10)
for (i in 1:10) {
  temp_model <- randomForest(Status ~ ., data = new_train, mtry = (i + 5))
  oob_values[i] <- temp_model$err.rate[nrow(temp_model$err.rate), 1]
}
ggplot(tibble("Error" = oob_values), aes(x = 6:15, y = Error)) +
  geom_line(aes(group = 1)) +
  labs(x = "Number of Variables",
       title = "Error VS mtry")
```

```
# choose mtry = 12 as the best model
best_model <- randomForest(Status ~ ., data = new_train, mtry = 12, ntree = 500)

valid_pred_rf <- valid %>%
  add_predictions(best_model)
table(valid_pred_rf$Status)
```

```
##
## Not Severe      Severe
##      42154      16077
```

```
confusionMatrix(valid_pred_rf$Status, valid_pred_rf$pred)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  Not Severe Severe
## Not Severe    33096   9058
## Severe        2143  13934
##
##              Accuracy : 0.8076
##              95% CI : (0.8044, 0.8108)
## No Information Rate : 0.6052
## P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.5753
##
```

```
## McNemar's Test P-Value : < 2.2e-16
##
##      Sensitivity : 0.9392
##      Specificity : 0.6060
##      Pos Pred Value : 0.7851
##      Neg Pred Value : 0.8667
##      Prevalence : 0.6052
##      Detection Rate : 0.5684
##      Detection Prevalence : 0.7239
##      Balanced Accuracy : 0.7726
##
##      'Positive' Class : Not Severe
##
```

(3) Decision tree

```
model_decision <- rpart(Status ~ ., data = new_train, method = "class")

valid_pred_dc <- valid %>%
  mutate(pred = predict(model_decision, valid, type = "class"))
confusionMatrix(table(valid_pred_dc$Status, valid_pred_dc$pred))
```

```
## Confusion Matrix and Statistics
##
##
##      Not Severe Severe
## Not Severe    30795 11359
## Severe        3894 12183
##
##      Accuracy : 0.7381
##      95% CI : (0.7345, 0.7416)
##      No Information Rate : 0.5957
##      P-Value [Acc > NIR] : < 2.2e-16
##
##      Kappa : 0.427
##
## McNemar's Test P-Value : < 2.2e-16
##
##      Sensitivity : 0.8877
##      Specificity : 0.5175
##      Pos Pred Value : 0.7305
##      Neg Pred Value : 0.7578
##      Prevalence : 0.5957
##      Detection Rate : 0.5288
##      Detection Prevalence : 0.7239
##      Balanced Accuracy : 0.7026
##
##      'Positive' Class : Not Severe
##
```